

# SAGE-CODE: Code-Augmented Reasoning on Adaptive Graphs

---

## A Comprehensive Technical Report

---

**Version:** 1.0.0

**Date:** January 2026

**Domain:** Clinical Intelligence & Data Quality Analytics

---

## Executive Summary

---

SAGE-CODE (**S**emantic **A**daptive **G**raph **E**ngine with **CODE** augmentation) is a novel hybrid Graph RAG (Retrieval-Augmented Generation) framework designed for clinical trial data intelligence. It combines knowledge graph traversal with dynamic code execution to provide actionable, data-backed insights for clinical trial management and data quality assurance.

The framework addresses the limitations of traditional RAG systems by:

1. **Adaptive Graph Traversal:** Using LLM-guided multi-hop reasoning to navigate complex clinical trial knowledge graphs
  2. **Code Augmentation:** Enabling on-the-fly Python code generation and execution for analytical queries
  3. **Semantic Selection:** Employing batched LLM scoring for intelligent candidate pruning during graph exploration
  4. **Chain-of-Thought Reasoning:** Integrating structured reasoning patterns for decision-making during traversal
- 

## Table of Contents

---

1. Introduction
2. Architecture Overview
3. Core Components

4. Knowledge Graph Construction
  5. Retrieval-Reasoning Pipeline
  6. Code Augmentation System
  7. Configuration & Customization
  8. Tools & Extensions
  9. Use Cases & Applications
  10. Technical Specifications
- 

# 1. Introduction

---

## 1.1 Problem Statement

Clinical trial management generates vast amounts of structured and semi-structured data across multiple domains:

- **Safety Events (eSAE)**: Serious Adverse Events requiring regulatory reporting
- **Data Discrepancies (EDRR)**: Open issues and data quality problems
- **Medical Coding (MedDRA/WHODD)**: Standardized coding for adverse events and medications
- **Visit Projections**: Patient visit scheduling and compliance tracking
- **Missing Data**: Form completion gaps and data entry delays

Traditional analytics approaches struggle with:

- Complex entity relationships spanning multiple data sources
- Need for both relational queries and analytical computations
- Requirement for explainable, data-backed decision support
- Integration of natural language queries with structured data analysis

## 1.2 SAGE-CODE Solution

SAGE-CODE addresses these challenges through a hybrid architecture that:

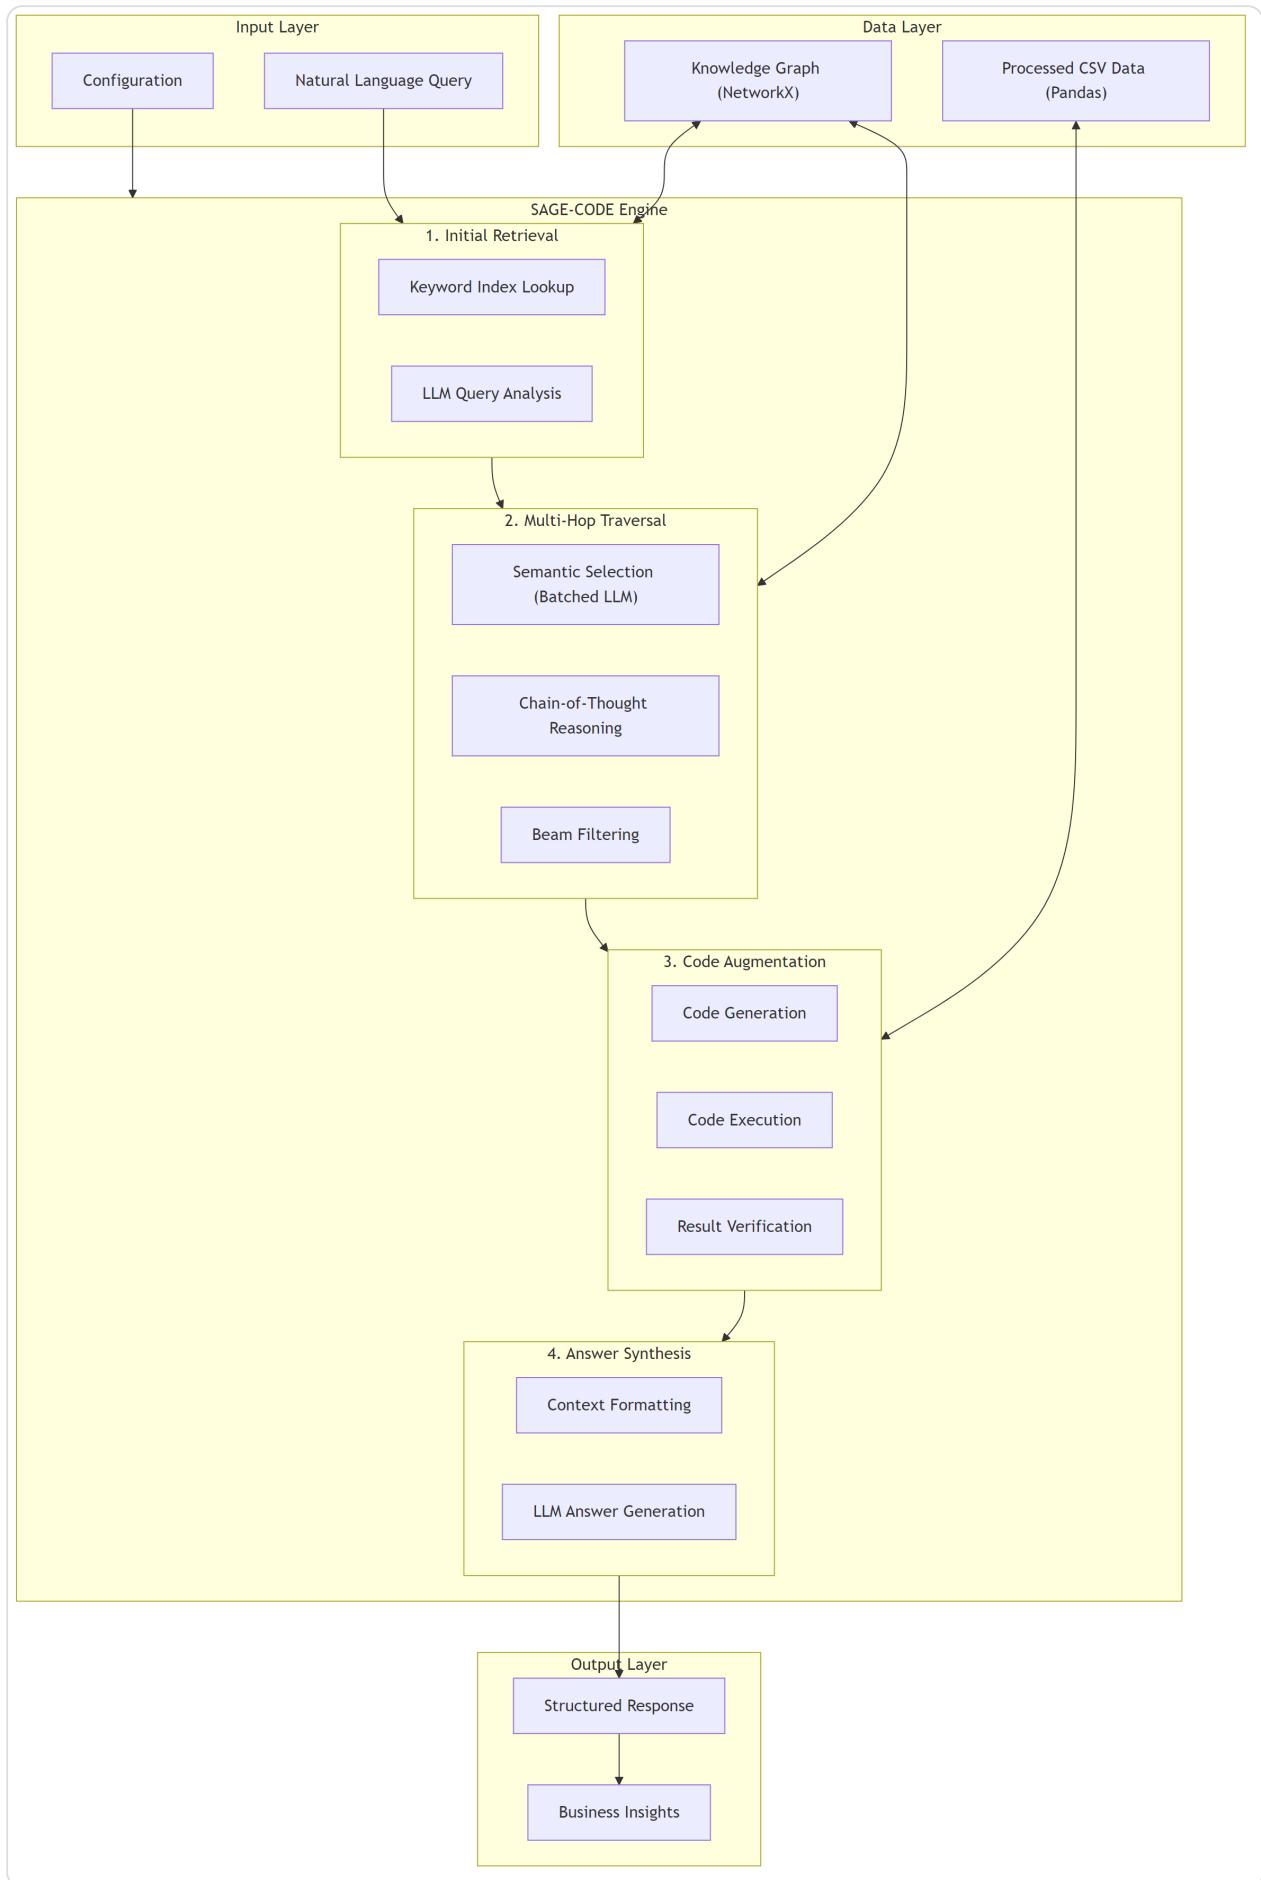
1. **Unifies Data**: Constructs a comprehensive knowledge graph from disparate CSV data sources
2. **Enables Natural Language Queries**: Allows clinical teams to ask questions in plain English
3. **Provides Grounded Answers**: Generates responses backed by actual data, not hallucinations
4. **Supports Analytics**: Executes Python code for complex aggregations and statistics

5. **Explains Reasoning:** Uses Chain-of-Thought patterns to show how conclusions were reached
-

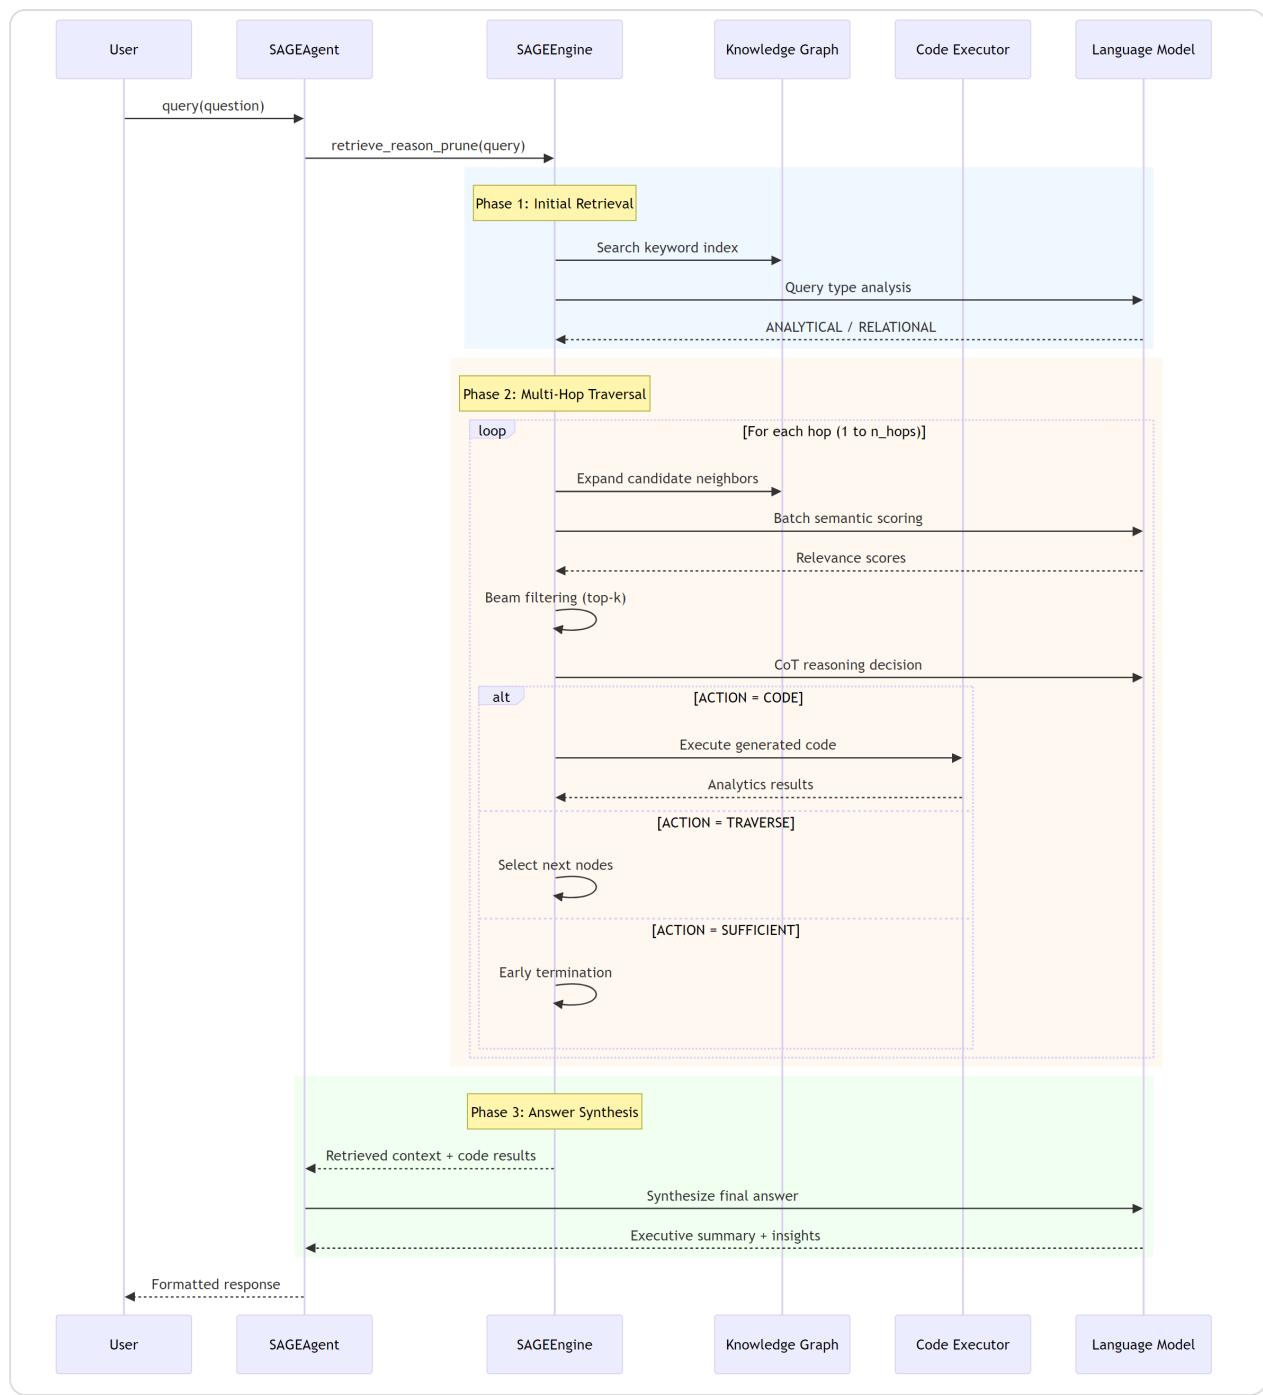
## **2. Architecture Overview**

---

### **2.1 High-Level System Architecture**

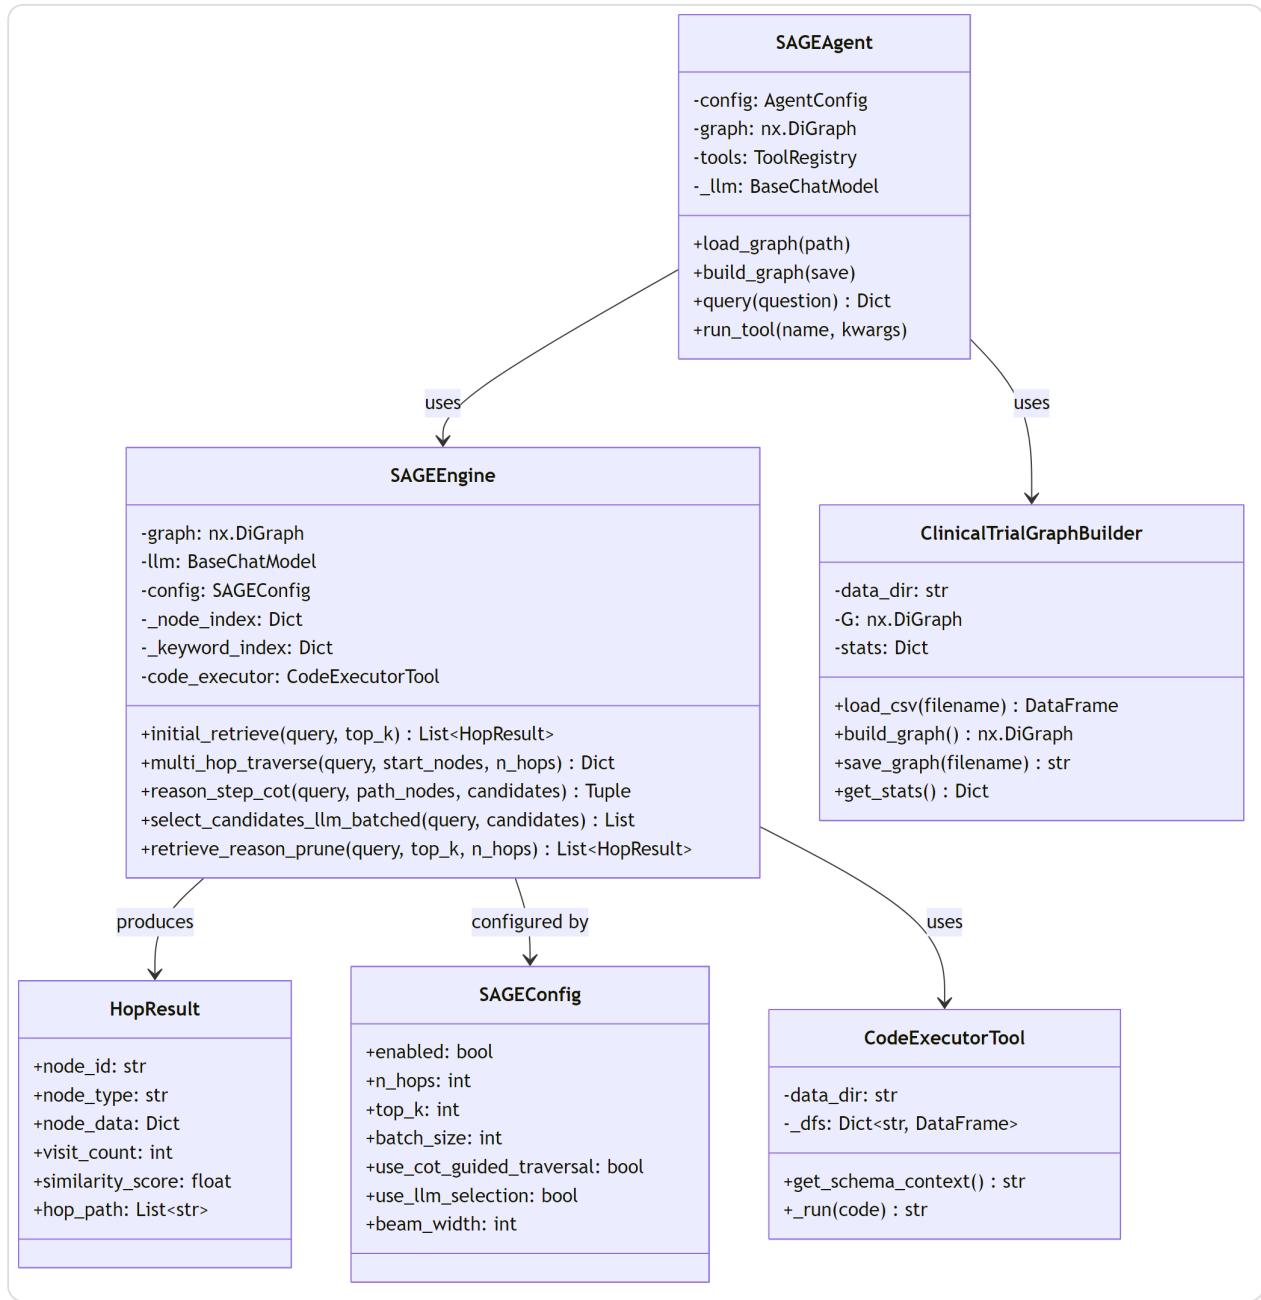


## 2.2 Component Interaction Flow



## 3. Core Components

### 3.1 Component Architecture



### 3.2 SAGEAgent

The `SAGEAgent` class serves as the primary interface for interacting with the SAGE-CODE system:

#### Key Responsibilities:

- Loading or building the knowledge graph
- Managing the tool registry

- Orchestrating queries through the SAGE engine
- Synthesizing final answers using LLM

#### **Initialization Flow:**

1. Load configuration (LLM settings, graph paths, SAGE parameters)
2. Load or build knowledge graph from CSV data
3. Initialize tool registry with graph and code execution tools
4. Set up LLM connection (supports Groq, Google, OpenAI providers)

### **3.3 SAGEEngine**

The `SAGEEngine` is the core retrieval-reasoning component:

#### **Key Features:**

- **Inverted Keyword Index:** O(1) lookup for initial node retrieval
- **LLM Query Analysis:** Distinguishes ANALYTICAL vs RELATIONAL queries
- **Batched LLM Scoring:** Efficient semantic relevance scoring
- **CoT-Guided Traversal:** Intelligent path selection with reasoning
- **Code Augmentation:** Dynamic code generation and execution

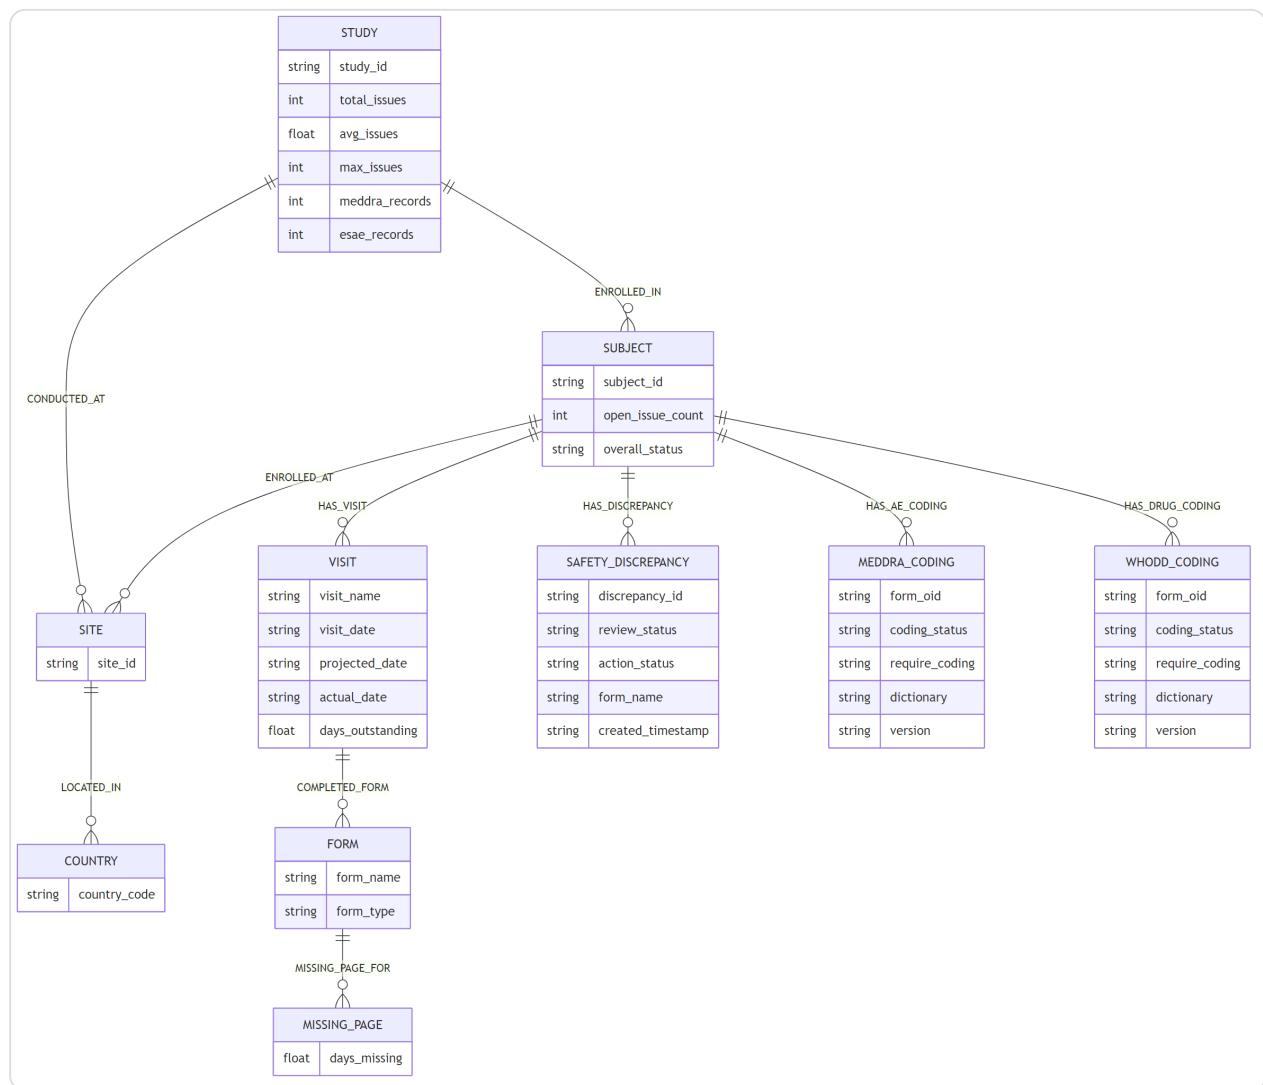
### **3.4 HopResult**

Data class representing a visited node during graph traversal:

```
@dataclass
class HopResult:
    node_id: str          # e.g., "SITE:637"
    node_type: str         # e.g., "Site", "Subject", "SafetyDiscrepancy"
    node_data: Dict        # All node attributes
    visit_count: int       # Times visited during traversal
    similarity_score: float # Relevance to query (0-1)
    hop_path: List[str]    # Path from initial node to this node
```

## 4. Knowledge Graph Construction

### 4.1 Graph Schema

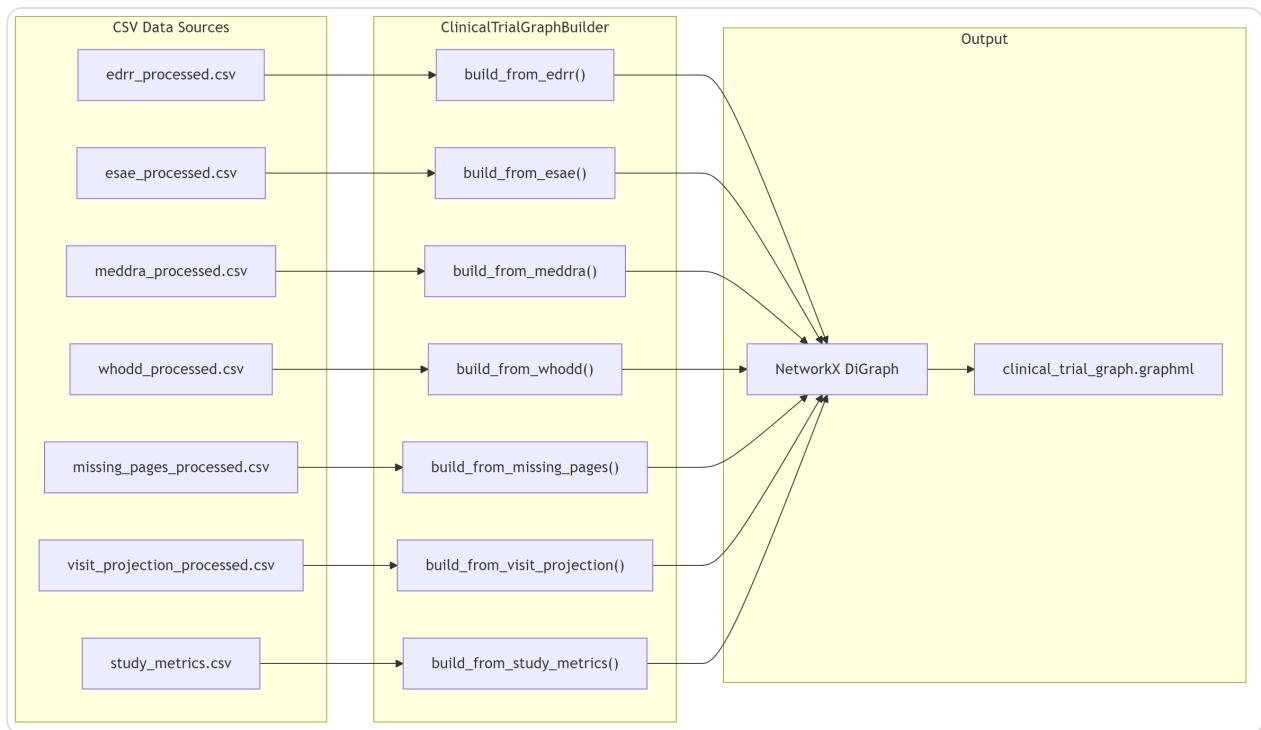


### 4.2 Data Sources

The `ClinicalTrialGraphBuilder` processes the following CSV files:

| Source File                    | Node Types Created  | Edge Types Created   |
|--------------------------------|---|--|
| edrr_processed.csv             | Study, Subject  | ENROLLED_IN  |
| esae_processed.csv             | Study, Country, Site,<br>Subject,<br>SafetyDiscrepancy        | LOCATED_IN, CONDUCTED_AT,<br>ENROLLED_IN, ENROLLED_AT,<br>HAS_DISCREPANCY                                |
| meddra_processed.csv           | Study, Subject,<br>MedDRACoding                               | ENROLLED_IN,<br>HAS_AE_CODING  |
| whodd_processed.csv            | Study, Subject,<br>WHODDCoding                                | ENROLLED_IN,<br>HAS_DRUG_CODING  |
| missing_pages_processed.csv    | Study, Country, Site,<br>Subject, Visit, Form,<br>MissingPage | LOCATED_IN, CONDUCTED_AT,<br>ENROLLED_IN, ENROLLED_AT,<br>HAS_VISIT, COMPLETED_FORM,<br>MISSING_PAGE_FOR |
| visit_projection_processed.csv | Study, Country, Site,<br>Subject, Visit                       | LOCATED_IN, CONDUCTED_AT,<br>ENROLLED_IN, ENROLLED_AT,<br>HAS_VISIT                                      |
| study_metrics.csv              | (Enriches existing<br>Study nodes)                            | -  |

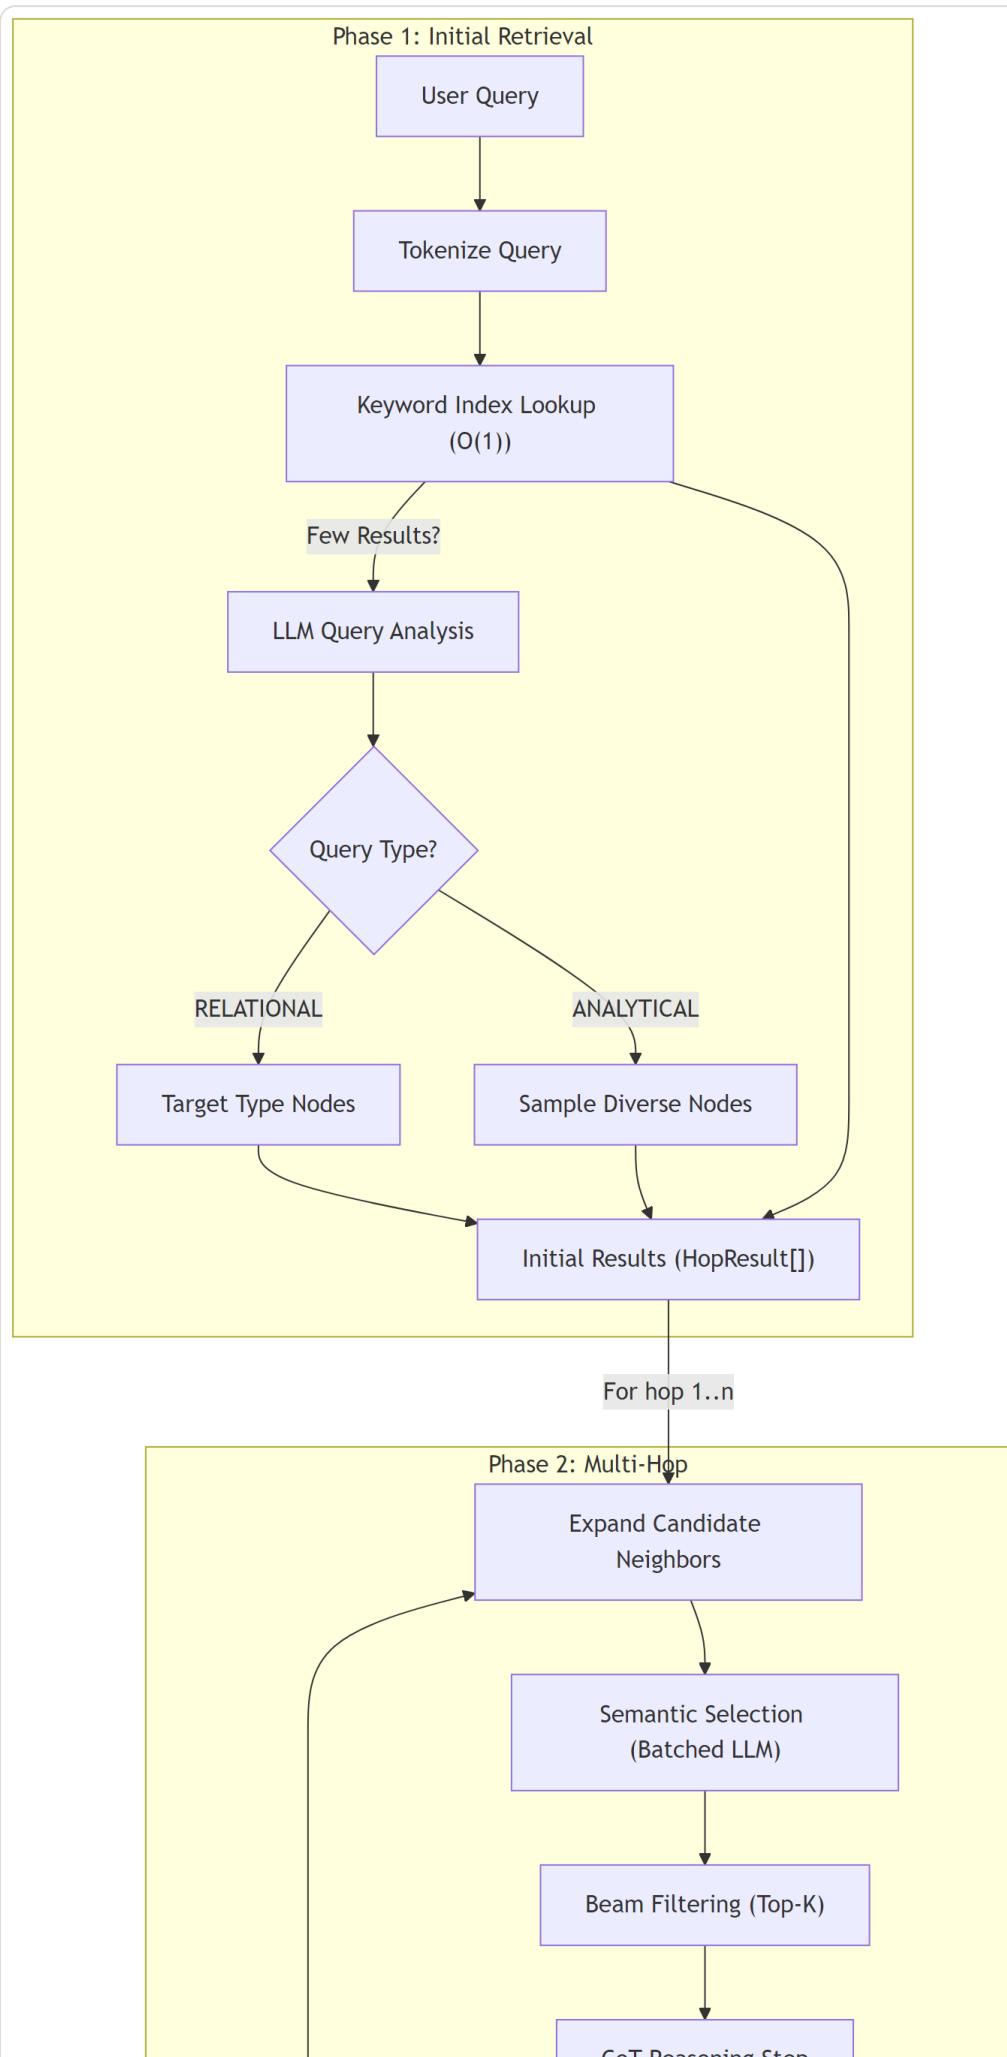
## 4.3 Graph Construction Process



## 5. Retrieval-Reasoning Pipeline

### 5.1 Pipeline Overview

The SAGE-CODE pipeline follows a **Retrieve** → **Reason** → **Prune** pattern:



## 5.2 Initial Retrieval

### Step 1: Keyword Index Lookup

- Query is tokenized into lowercase words
- Inverted index provides O(1) lookup: `word → [node_id, ...]`
- Nodes are scored by match overlap: `score = matches / query_words`

### Step 2: LLM Query Analysis (if insufficient matches)

```
analysis_prompt = """
Analyze this clinical trial query and determine:
1. query_type: "ANALYTICAL" or "RELATIONAL"
2. target_types: List of node types to search
3. key_entities: Specific entity IDs mentioned
"""
```

### Step 3: Query Type Handling

- **ANALYTICAL**: Sample diverse nodes for code context
- **RELATIONAL**: Target specific node types for traversal

## 5.3 Semantic Selection (Batched)

The `select_candidates_llm_batched` function scores candidates using LLM:

```
BATCH_SELECTION_PROMPT = """
Rate relevance (0-10) of each candidate node for the query.
Consider:
1. Explicit matches (keywords)
2. Semantic relationships
3. DATA POTENTIAL: Score HIGHER if the node unlocks computation

Query: "{query}"
Candidates: {candidates_text}

Response: {"scores": {"candidate_index_0": 9, ...}}
"""
```

## 5.4 Chain-of-Thought Reasoning

The CoT step uses a code-aware prompt that provides:

- Current path context (visited nodes)
- Candidate nodes for traversal
- Available DataFrame schemas

### **Decision Actions:**

| Action     | Description                      | Effect                               |
|------------|----------------------------------|--------------------------------------|
| CODE       | Generate and execute Python code | Analytics query, may terminate early |
| TRAVERSE   | Select nodes for next hop        | Continue graph exploration           |
| SUFFICIENT | Information complete             | Terminate traversal                  |

## **5.5 Helpfulness Scoring**

Final results are scored using:

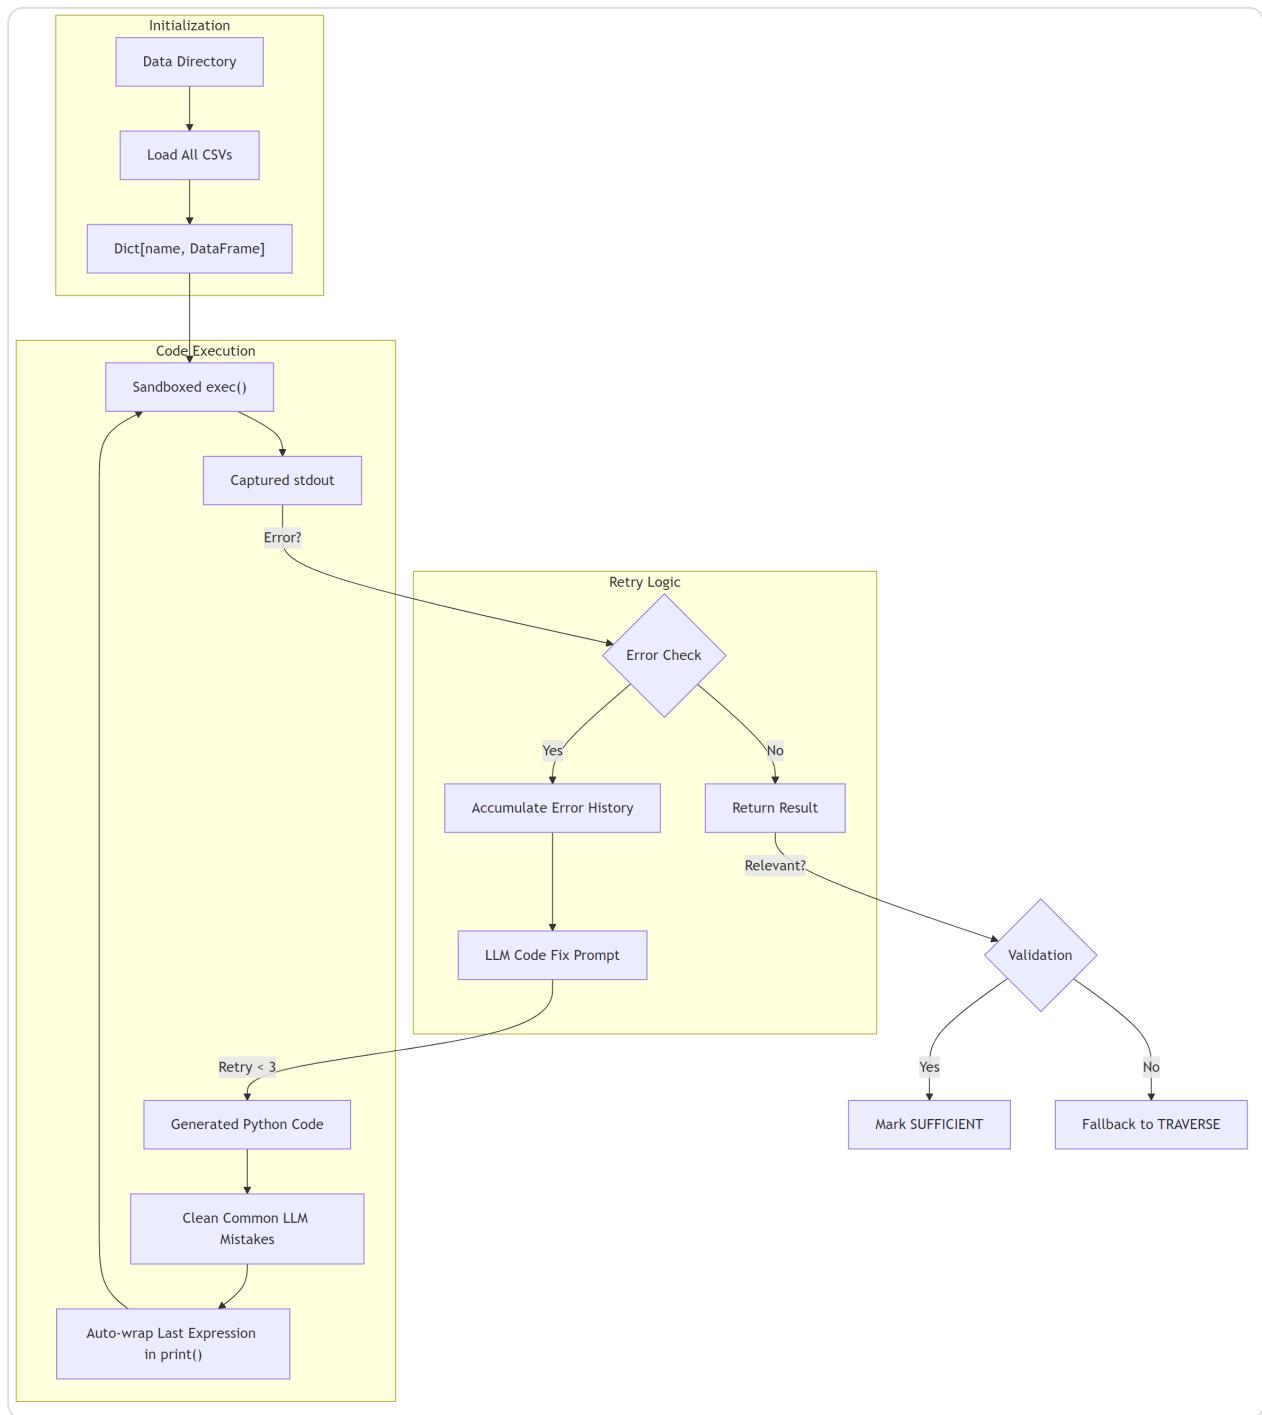
```
helpfulness = similarity_weight * sim_score + (1 - similarity_weight) * importance_score
```

Where:

- `sim_score` : Semantic similarity to query
  - `importance_score` : Normalized visit count during traversal
-

# 6. Code Augmentation System

## 6.1 Code Executor Architecture



## 6.2 Available DataFrames

The code executor dynamically loads all CSVs from the processed data directory:

| DataFrame Name                | Source File                    | Key Columns  |
|-------------------------------|--------------------------------|--|
| esae_processed_df             | esae_processed.csv             | study_id, site, patient_id, review_status                  |
| missing_pages_processed_df    | missing_pages_processed.csv    | study_name, sitenumber, subjectname, no__days_page_missing |
| meddra_processed_df           | meddra_processed.csv           | study, subject, coding_status                              |
| whodd_processed_df            | whodd_processed.csv            | study, subject, coding_status                              |
| visit_projection_processed_df | visit_projection_processed.csv | study, site, subject, __days_outstanding                   |
| study_metrics_df              | study_metrics.csv              | study, total_issues, avg_issues                            |
| edrr_processed_df             | edrr_processed.csv             | study, subject, total_open_issue_count_per_subject         |

## 6.3 Code Generation Example

**Query:** "Which sites have the most pending safety reviews?"

**Generated Code:**

```
result = esae_processed_df[
    esae_processed_df['review_status'] == 'Pending for Review'
].groupby('site').size().sort_values(ascending=False).head(10)
print(result)
```

**Output:**

```
site
637      45
412      38
...
...
```

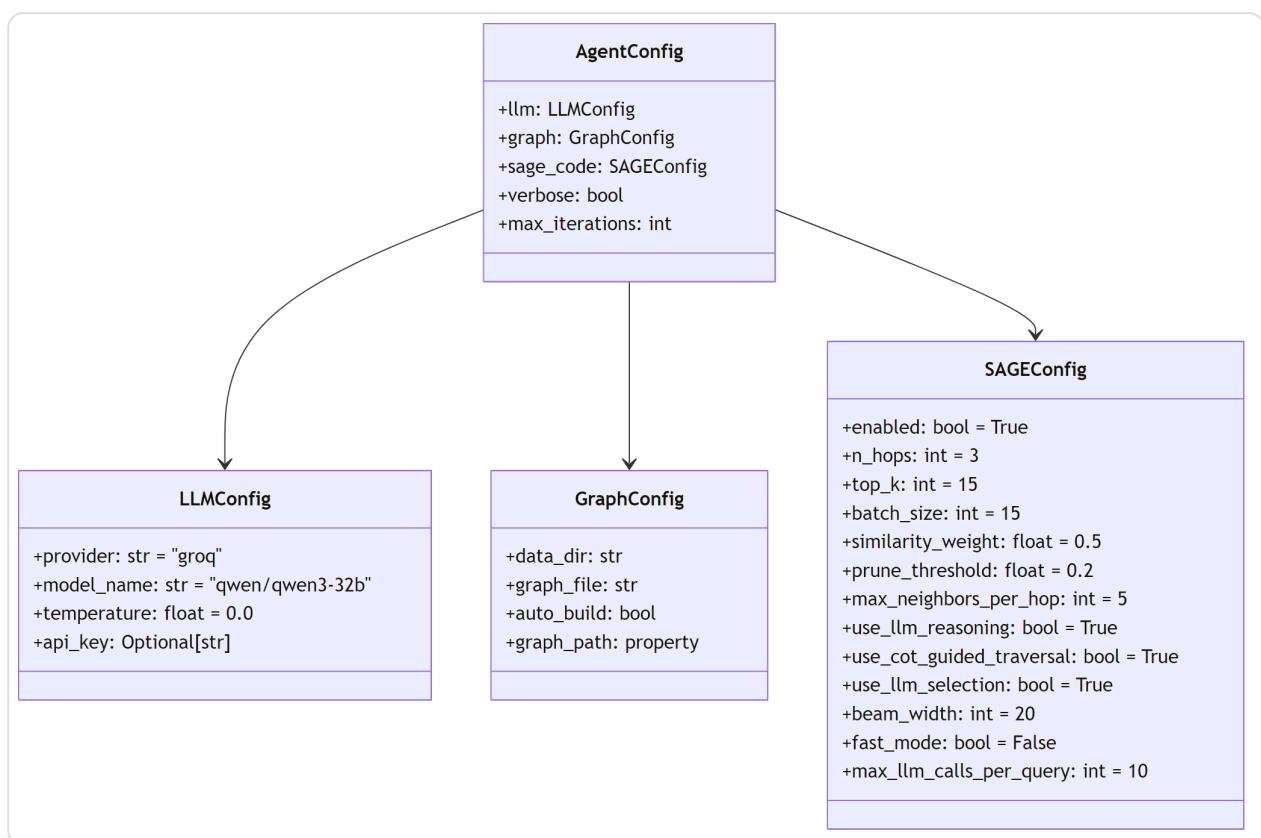
## 6.4 Self-Healing Code Execution

The system implements automatic error recovery:

1. **Error Detection:** Check if output starts with "Error:"
2. **Error Accumulation:** Track all errors across retries
3. **LLM Fix Request:** Provide full error history to avoid repeated mistakes
4. **Maximum Retries:** 3 attempts before falling back to traversal

## 7. Configuration & Customization

### 7.1 Configuration Classes



## 7.2 Key Configuration Parameters

| Parameter                | Default | Description                               |
|--------------------------|---------|---|
| n_hops                   | 3       | Maximum traversal depth                   |
| top_k                    | 15      | Number of results to return               |
| batch_size               | 15      | LLM batch size for scoring                |
| beam_width               | 20      | Candidates kept per hop                   |
| similarity_weight        | 0.5     | Weight for semantic vs structural scoring |
| prune_threshold          | 0.2     | Minimum helpfulness score                 |
| use_cot_guided_traversal | True    | Enable CoT reasoning                      |
| use_llm_selection        | True    | Enable semantic scoring                   |
| fast_mode                | False   | Skip LLM reasoning (keyword only)         |
| skip_multi_hop           | False   | Return only initial retrieval             |

## 7.3 Performance Modes

### Full Mode (Default):

- All LLM features enabled
- Best accuracy for complex queries
- Higher latency and API cost

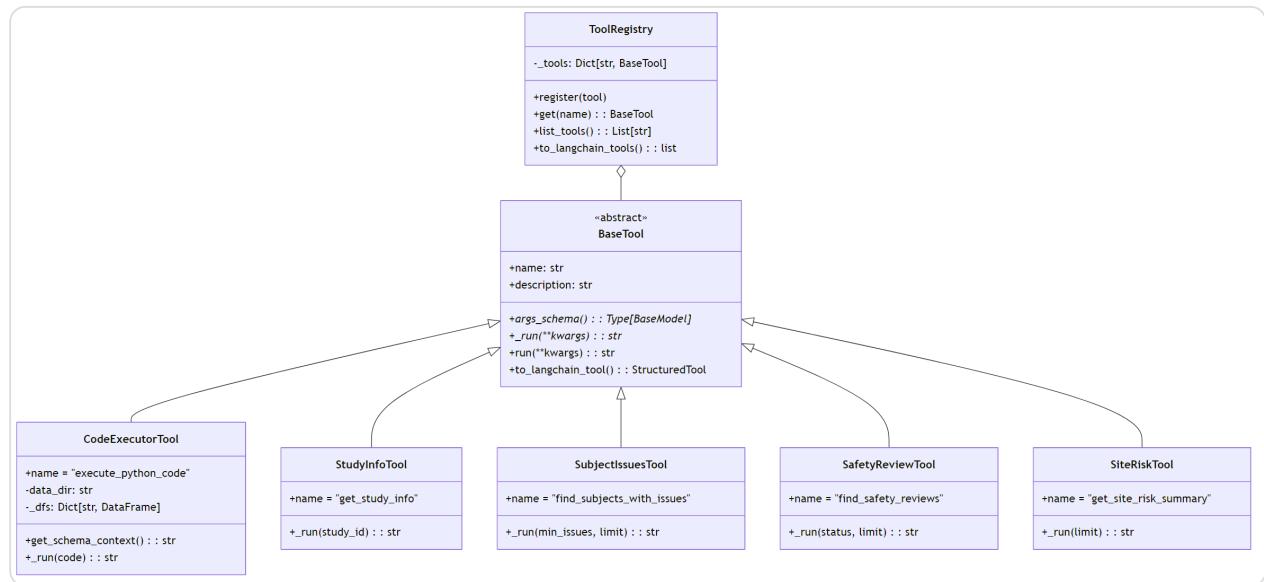
### Fast Mode:

```
config = SAGEConfig(  
    fast_mode=True,  
    skip_multi_hop=True,  
    use_llm_reasoning=False  
)
```

- Keyword retrieval only
- Minimal latency
- Suitable for simple lookups

# 8. Tools & Extensions

## 8.1 Tool Architecture



## 8.2 Available Tools

| Tool Name                               | Description                         | Key Parameters                            |
|---|-------------------------------------|---|
| <code>execute_python_code</code>        | Run pandas code on clinical data    | <code>code: str</code>                    |
| <code>get_study_info</code>             | Get study information and metrics   | <code>study_id: str</code>                |
| <code>find_subjects_with_issues</code>  | Find subjects with open issues      | <code>min_issues: int , limit: int</code> |
| <code>find_safety_reviews</code>        | Find safety discrepancies by status | <code>status: str , limit: int</code>     |
| <code>get_safety_reviews_by_site</code> | Aggregate safety reviews by site    | <code>status: str , limit: int</code>     |
| <code>find_missing_pages</code>         | Find aged missing pages             | <code>min_days: float , limit: int</code> |
| <code>get_site_risk_summary</code>      | Get risk summary by site            | <code>limit: int</code>                   |

## 8.3 DQI Analytics Tools (Extension)

Additional analytics capabilities:

| Tool                                  | Description                        |
|---------------------------------------|------------------------------------|
| <code>calculate_dqi</code>            | Calculate Data Quality Index score |
| <code>explain_dqi</code>              | LLM-generated DQI analysis         |
| <code>rank_sites</code>               | Rank sites by metrics              |
| <code>benchmark_site</code>           | Compare site against peers         |
| <code>identify_underperformers</code> | Find sites needing attention       |

## 9. Use Cases & Applications

---

### 9.1 Clinical Trial Data Quality

**Example Query:** "Is the current data snapshot clean enough for interim analysis?"

**SAGE-CODE Response:**

```
### Executive Summary
The data shows moderate readiness with 23% of sites having pending safety reviews.

### Key Findings
| Metric | Value | Status |
|-----|-----|-----|
| Open Issues | 1,247 | ⚠️ Elevated |
| Pending Reviews | 456 | Action Required |
| Missing Pages >30d | 89 | Needs Resolution |

### Recommendations
1. Prioritize Site 637 for data cleanup (45 pending reviews)
2. Resolve missing pages before lock
3. Complete MedDRA coding for 12 outstanding AEs
```

### 9.2 Site Performance Analysis

**Example Query:** "Which sites have the most data quality issues?"

**Process:**

1. CoT decides: ACTION = CODE
2. Generates aggregation code
3. Executes against esae\_processed\_df
4. Returns ranked site list with issue counts

### 9.3 Subject Safety Tracking

**Example Query:** "Show me subjects with more than 5 open issues enrolled at Site 412"

**Process:**

1. Initial retrieval finds Site 412 node
2. Traversal explores ENROLLED\_AT edges
3. Filters subjects by issue count
4. Returns subject list with details

# 10. Technical Specifications

---

## 10.1 Dependencies

| Package                | Version | Purpose                |
|------------------------|---------|------------------------|
| networkx               | ≥2.8    | Graph data structure   |
| pandas                 | ≥1.5    | DataFrame operations   |
| pydantic               | ≥2.0    | Data validation        |
| langchain-core         | ≥0.1    | LLM abstraction        |
| langchain-groq         | ≥0.1    | Groq LLM integration   |
| langchain-google-genai | ≥0.1    | Google LLM integration |
| langchain-openai       | ≥0.1    | OpenAI LLM integration |

## 10.2 LLM Providers

| Provider | Model          | Use Case                   |
|----------|----------------|----------------------------|
| Groq     | qwen/qwen3-32b | Default, fast inference    |
| Google   | gemini-1.5-pro | Multi-modal, large context |
| OpenAI   | gpt-4o         | High accuracy              |

## 10.3 Performance Characteristics

| Metric               | Typical Value |
|----------------------|---------------|
| Initial Retrieval    | <100ms        |
| Per-Hop LLM Call     | 500-2000ms    |
| Code Execution       | 50-500ms      |
| Total Query (3 hops) | 3-10 seconds  |
| LLM Calls per Query  | 3-12          |

## 10.4 Graph Statistics (Typical)

| Metric      | Value    |
|-------------|----------|
| Total Nodes | 200,000+ |
| Total Edges | 500,000+ |
| Node Types  | 10+      |
| Edge Types  | 12+      |

---

## Appendix A: LLM Prompts

---

### A.1 SAGE Agent System Prompt

You are a Senior Clinical Trial Consultant AI. Your role is to provide clear, actionable business insights derived from complex data.

## CRITICAL RESPONSE GUIDELINES

1. Human-Centric & Professional: Write for business stakeholders
2. No Technical Details in Output
3. Data-Backed Insights: Always include specific numbers
4. Action-Oriented: Conclude with specific next steps

## A.2 Code-Augmented CoT Prompt

You are an intelligent clinical trial data analyst with Python Pandas capabilities.

Goal: Answer the user's query using the available dataframes.

PREFER CODE over traversal for analytical questions.

DECISION RULES:

1. Use CODE if the query asks for: counts, aggregations, rankings, statistics
2. Use TRAVERSE only if you need to find specific entity IDs or relationships
3. Use SUFFICIENT only if the Current Knowledge already contains the answer

## A.3 Batch Selection Prompt

Rate relevance (0-10) of each candidate node for the query.

Consider:

1. Explicit matches (keywords)
2. Semantic relationships
3. DATA POTENTIAL: Score HIGHER if the node unlocks computation

## Appendix B: File Structure

```
sage_code/
├── __init__.py           # Package exports
├── agent.py              # SAGEAgent class
├── engine.py             # SAGEEngine (core algorithm)
├── graph_builder.py      # ClinicalTrialGraphBuilder
├── config.py              # Configuration dataclasses
├── models.py              # HopResult dataclass
├── prompts.py             # LLM prompt templates
└── tools/
    ├── __init__.py        # Tools exports
    ├── base_tool.py        # BaseTool, ToolRegistry
    ├── code_executor.py    # CodeExecutorTool
    ├── graph_tools.py      # Graph query tools
    └── dqi_analytics_tools.py # DQI extension tools
```

## References

---

1. **Graph RAG:** Lewis, P. et al. (2020). Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks.
  2. **Chain-of-Thought Prompting:** Wei, J. et al. (2022). Chain-of-Thought Prompting Elicits Reasoning in Large Language Models.
  3. **Knowledge Graphs for Clinical Trials:** Various industry implementations.
- 

*Document generated for Novartis Clinical Intelligence Platform  
SAGE-CODE Framework v1.0.0*